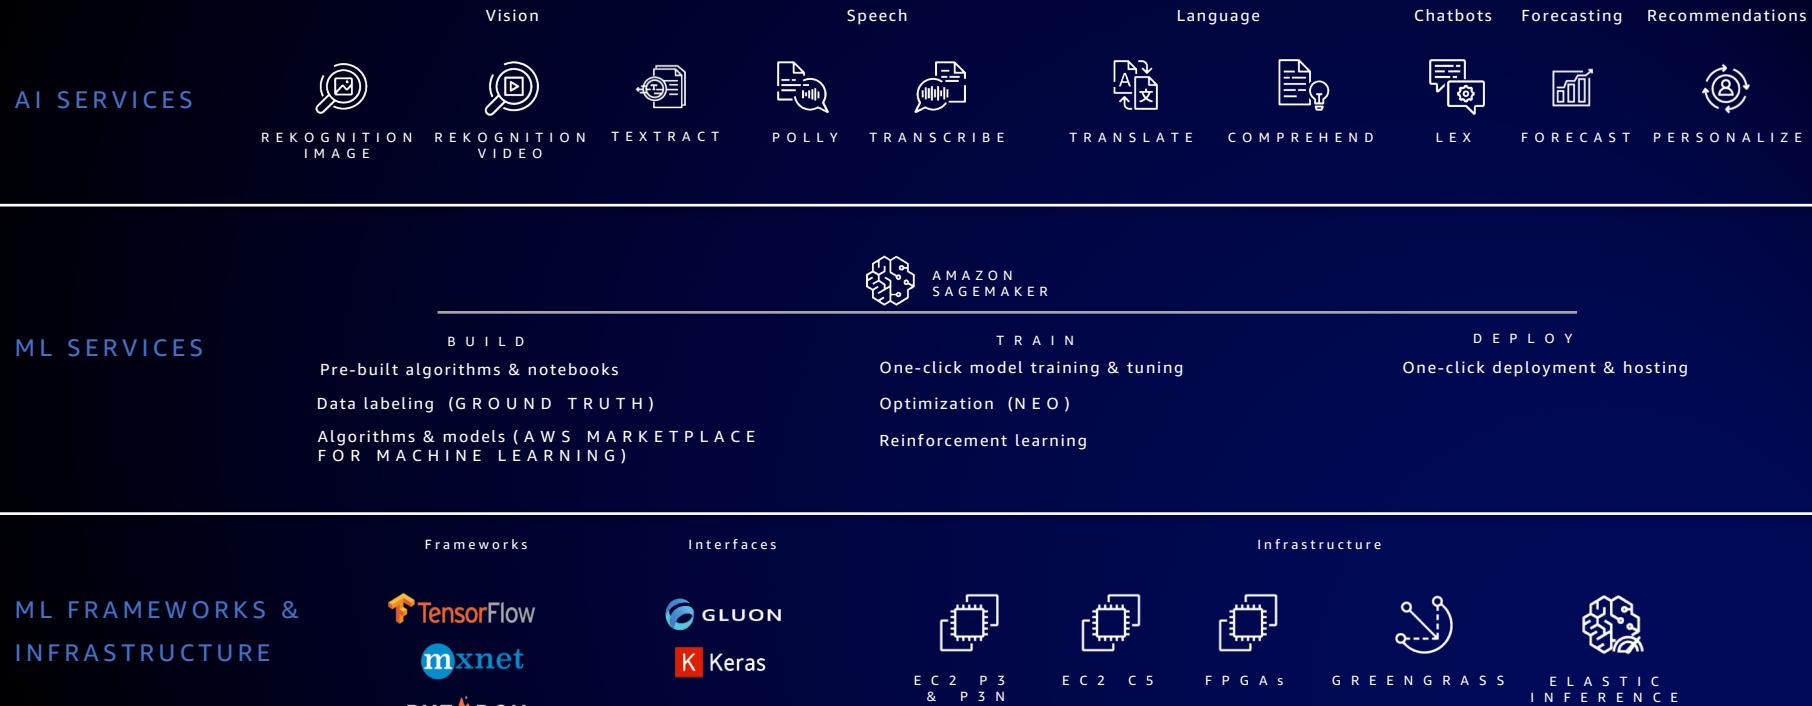


# Amazon ML Services

Pedro Paez

*Specialist Solutions Architect AI/ML  
Amazon Web Services*

# The Amazon ML stack: Broadest & deepest set of capabilities





# Amazon SageMaker

A **fully managed service** that enables **data scientists** and **developers** to quickly and easily **build** machine-learning based models **into production** smart applications.

# Amazon SageMaker: build, train, and deploy ML

Pre-built notebooks for common problems

Built-in, high performance algorithms

One-click training

Optimization

One-click deployment

Fully managed with auto-scaling

Choose your ML algorithm

Optimize your ML algorithm

Set up and manage environments for training

Train and tune model (trial and error)

Deploy model in production

Scale and manage the production environment



# Amazon SageMaker

Build

Pre-built  
notebook  
instances

Highly-optimized  
machine learning  
algorithms



# Amazon SageMaker

Build

Pre-built  
notebook  
instances

Highly-optimized  
machine learning  
algorithms

Easier training with  
hyperparameter  
optimization



Train



# Amazon SageMaker

Deploy

Build

Fully-managed hosting at scale

Deployment without engineering effort

Pre-built notebook instances

Highly-optimized machine learning algorithms

Easier training with hyperparameter optimization

One-click training for ML, DL, and custom algorithms



Train

TensorFlow

PYTORCH

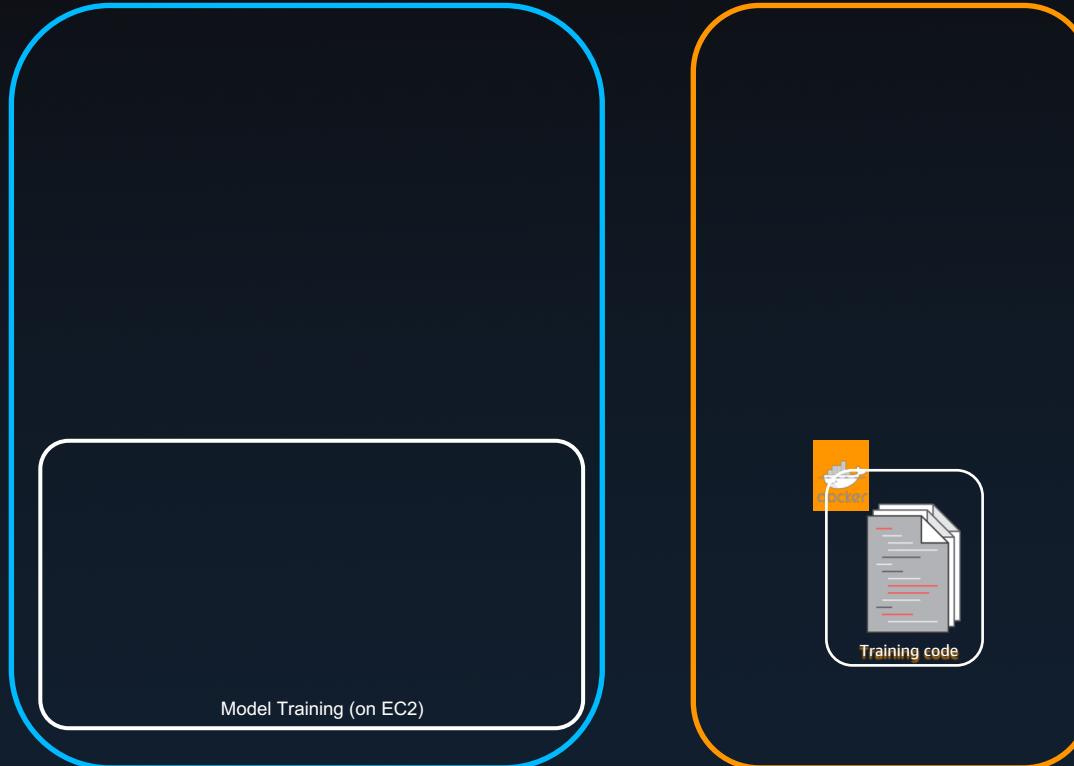
Client application

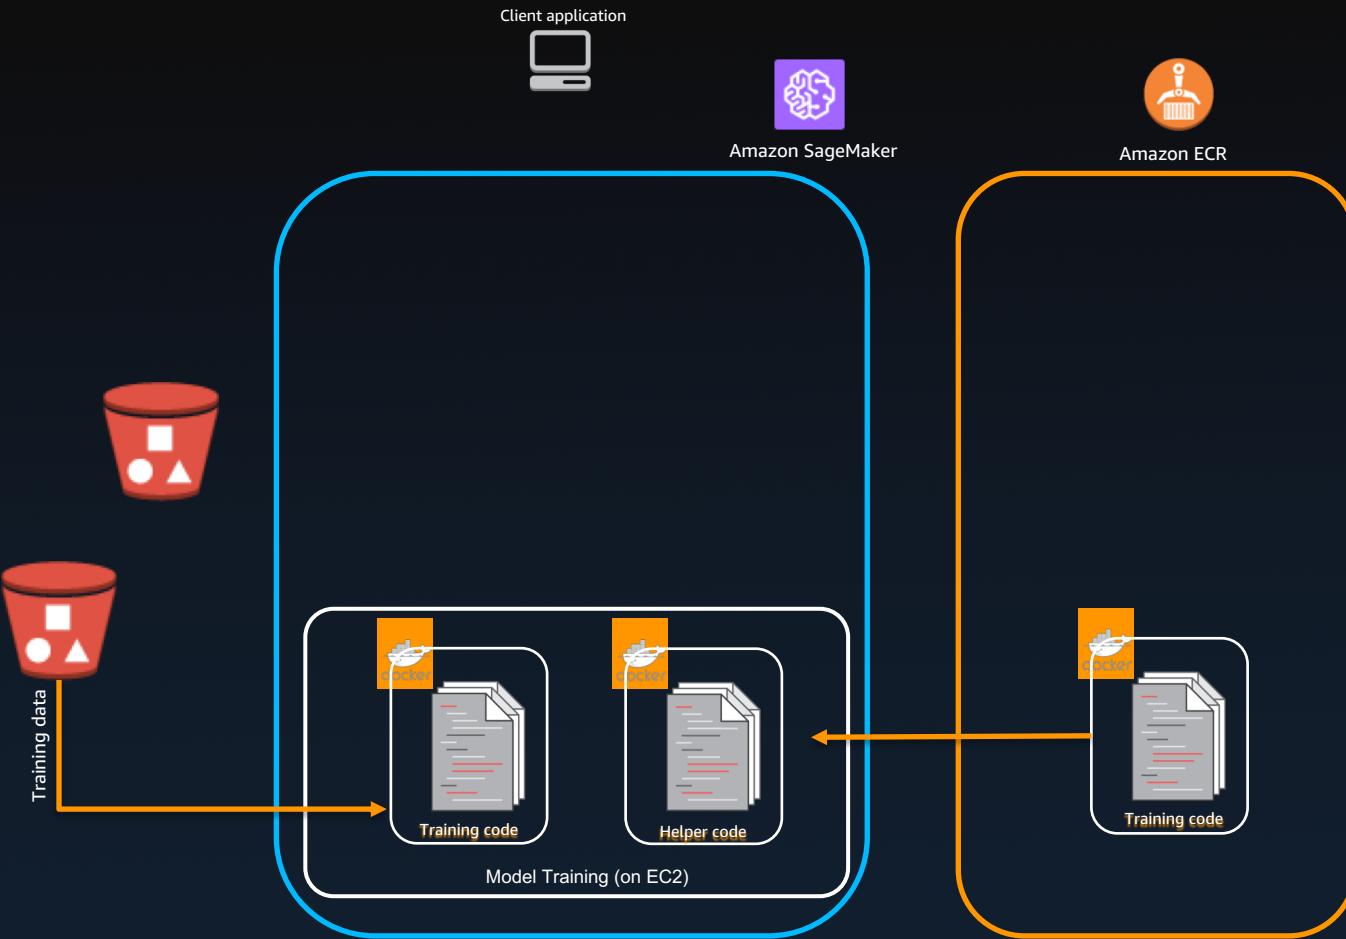


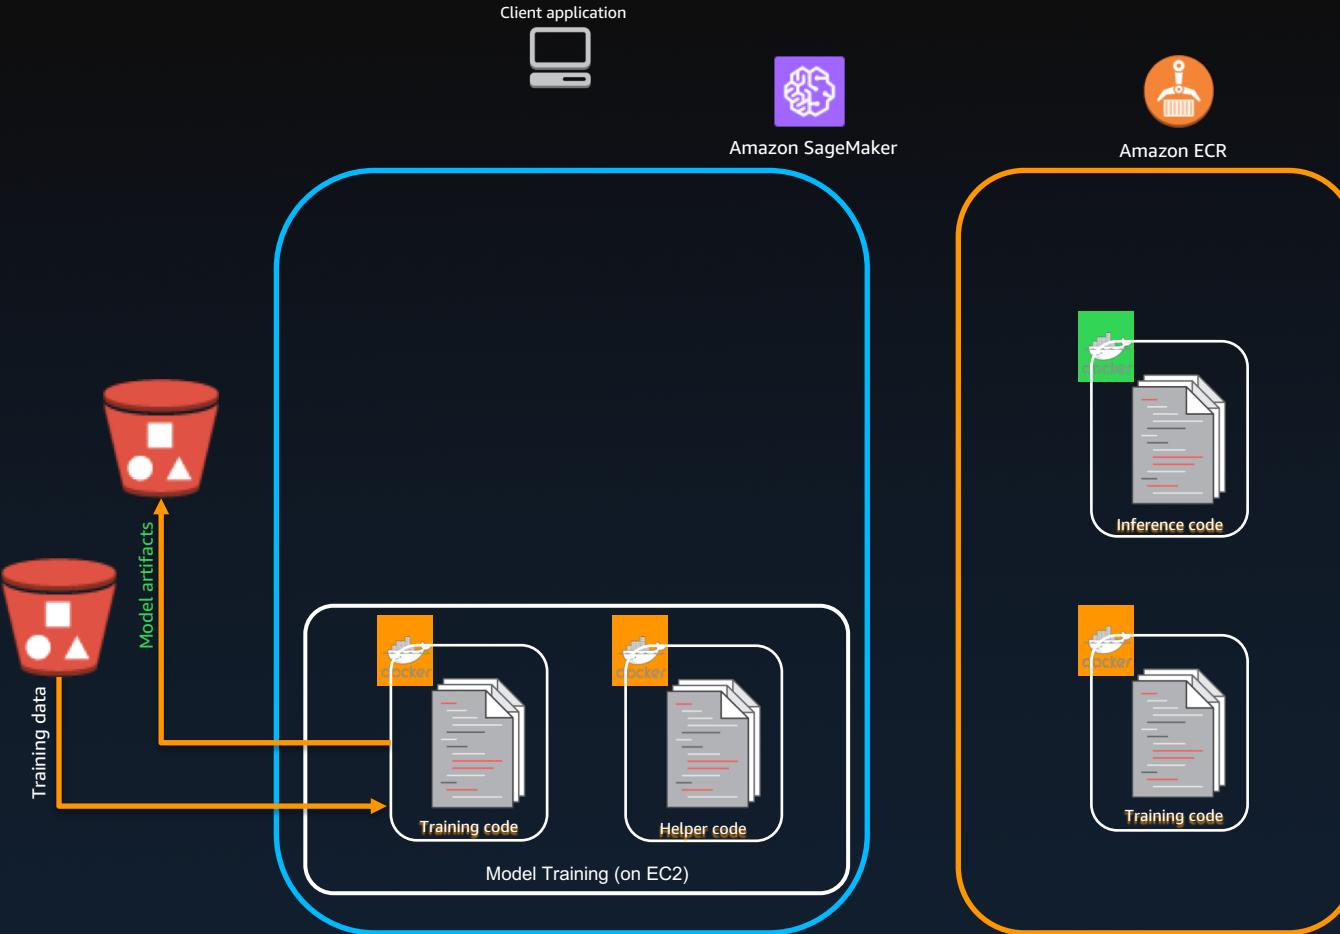
Amazon SageMaker

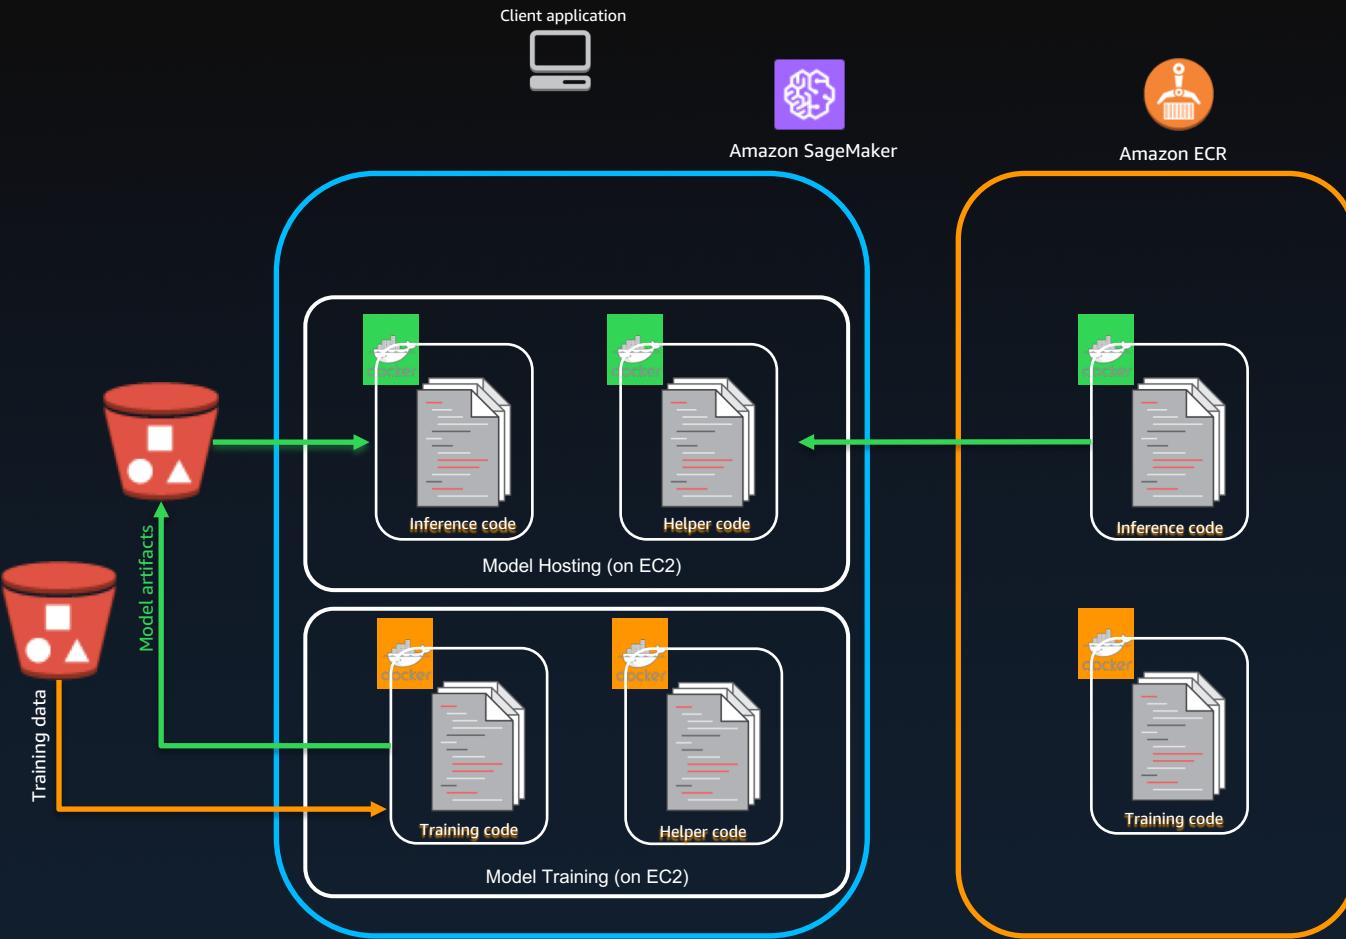


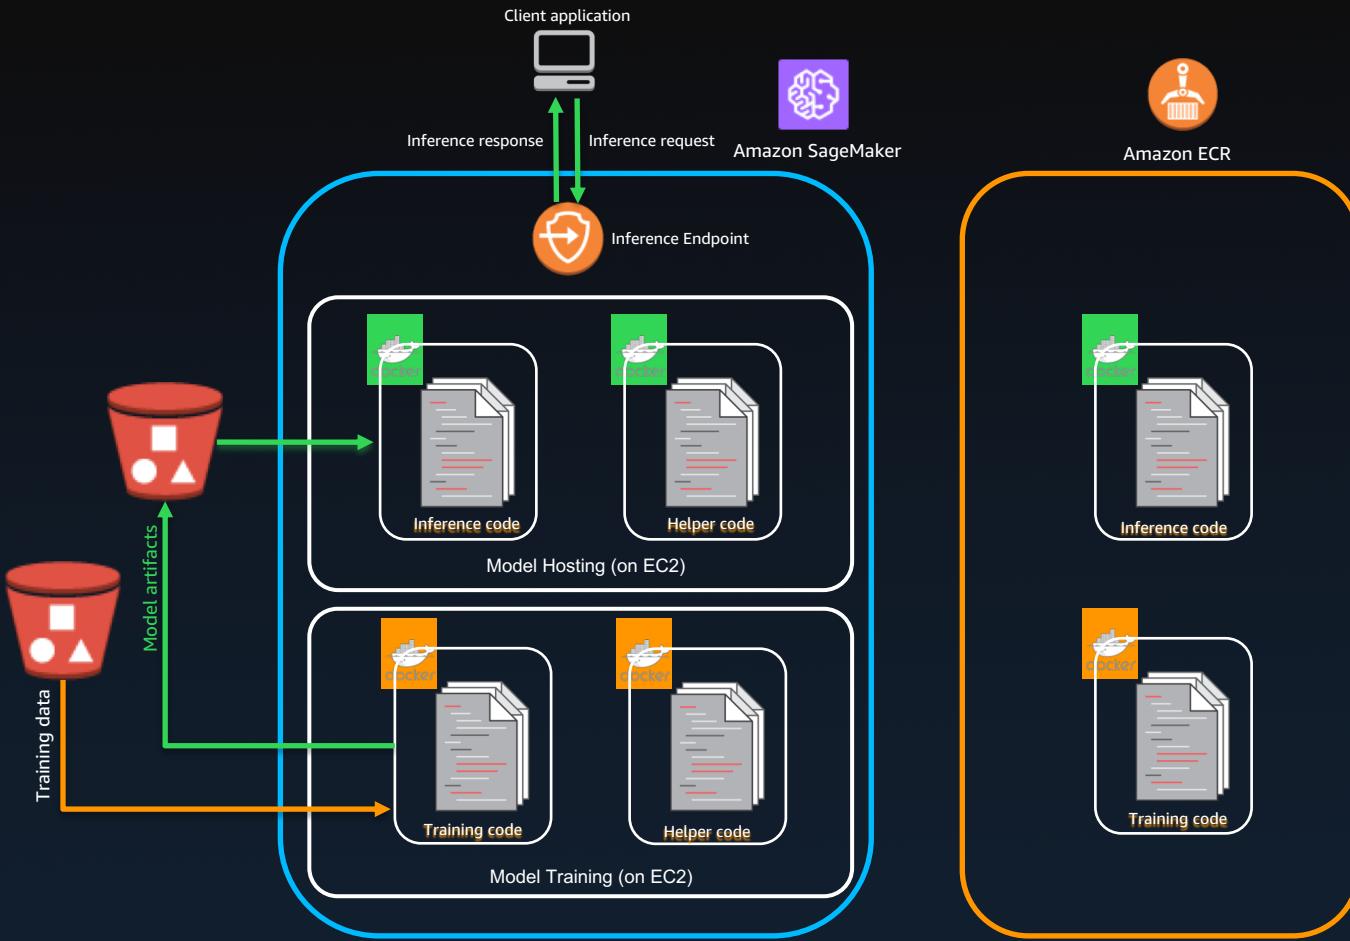
Amazon ECR











# Reference Architecture

